# TALENTRA DATABASE DESIGN HACKATHON

B19 — Jobs & Applications Design Decision Register

**Status:** Working proposal for team review. This document does not claim that B19 has passed or been signed off.

## 1. Where we are

The docket defines B18 as combining the approved Jobs & Applications entities and columns into a design draft. B19 validates that design. The design should not move to Milestone 6 until the team has reconciled entity names, relationships, columns, ownership boundaries, and unresolved business rules.

## 2. Decision register

| ID / topic | Recommended resolution (proposal) | Why it needs team confirmation |
|---|---|---|
| D1 — Canonical names | Use exactly the names in the approved Milestone 2 entity list. Align the ERD, column dictionary, and validation checklist to those names. Do not keep parallel names such as `interview` and `job_interview`. | The docket requires design to be based on the validated entity inventory; it does not prescribe every physical table name. |
| D2 — Application status | Keep current status on `application` for efficient reads and an append-only `application_status_history` for the lifecycle trail. Define allowed transitions and who can trigger them. | The docket names the pipeline (applied, shortlisted, interviewed, offered, joined, rejected), but not transition rules or history implementation. |
| D3 — Questionnaire answers | If candidates answer job-specific questions, use a response entity tied to the application and question; enforce one response per application/question unless multi-answer questions are required. | Questionnaire and question concepts are in scope; answer storage/cardinality must be agreed. |
| D4 — Joining / rejection | Represent joining as an explicit lifecycle event or application outcome; rejection should be a status/outcome with timestamp and optional reason. Avoid a second independent state that can contradict application status. | The docket names joins and rejections but does not dictate whether they are tables, events, or status values. |
| D5 — Offers | If offer details must be retained, model an offer record linked to an application. Decide whether an application may have multiple offer versions; if not, enforce one active offer and retain history separately if needed. | Offer cardinality and revision handling are not specified in the docket. |
| D6 — Reapplications | Permit multiple applications by the same candidate to the same job over time unless product rules prohibit it. Define the natural uniqueness rule (for example, one active application per candidate/job). | The docket does not specify whether repeat or withdrawn applications are allowed. |
| D7 — Shortlists | Use an application-linked shortlist entity or event if the system must track shortlist membership, recruiter, and time. Avoid a candidate-only bookmark table for recruiter shortlisting. | The required shortlist concept is in scope, but exact ownership and history are not defined. |
| D8 — Master Data references | Reference Master Data for job role, employment/work type, skills, language and location where those domains are owned there. Do not duplicate lookup definitions inside Jobs & Applications. | The docket assigns these shared domains to Master Data; exact FK targets must match its approved schema. |
| D9 — Cross-module entities | Show `candidate_profile`, organisation/account, location, and Master Data entities as dependencies where appropriate; mark Phase 2-owned entities as external references rather than fully designed Phase 1 tables. | The docket separates common Phase 1 modules from later modules. |
| D10 — Delete/update policy | For each FK, document ON UPDATE and ON DELETE behavior. Prefer restrictive deletion for referenced business history; define anonymisation/retention separately from cascading deletes. | The docket explicitly requires DB-enforced FKs and explicit update/delete rules, plus a candidate retention/deletion policy. |
| D11 — Published job edits | Decide whether editing a published requirement updates the same record or creates a version/snapshot. If candidates must see the original job terms, preserve a versioned snapshot. | Versioning behaviour is not specified; it affects auditability and application fairness. |
| D12 — Column completeness | For every entity, complete column, type, nullability rationale, classification, audit fields, deletion policy, and default. Keep all unresolved columns marked OPEN. | Milestone 4 draft itself notes that several inventories and standards still require completion/review. |

## 3. B19 exit criteria

• The approved entity list and the design use identical canonical names.

• Every relationship has cardinality, optionality, and a documented FK target.

• Every entity has a complete column inventory, including type, nullability and rationale, data classification, and audit fields.

• All required lifecycle concepts (publish state, application pipeline, shortlist, interview, offer, join/rejection, and relevant history) are represented consistently.

• Cross-module ownership and Master Data references are explicit; Phase 2 modules are not accidentally claimed as completed Phase 1 designs.

• Open decisions are either resolved by the team or visibly recorded as accepted assumptions with an owner.

• The team records B19 validation outcome and approver/date. Until then, status remains 'not signed off'.

## 4. Next practical action

Hold a short team review and resolve D1–D12. Then update the canonical entity list, ERD, and column dictionary in one pass; run the B19 checklist against that single source of truth. Only after B19 is accepted should the team start B20/B22 (keys and constraints).

**Important:** These are recommendations, not facts stated by the docket. The docket supplies the required scope and standards; the proposed resolutions above are design suggestions for the team to approve or reject.